

# **The Thiele Machine**

A Complete Mathematical Specification

Derived from the Fully Machine-Checked Coq Corpus

Formal Specification

## Abstract

This document provides a complete, standalone mathematical specification of the Thiele Machine, a formal model of computation in which structural information carries an explicit, conserved cost  $\mu$ . The document is organized so that the machine can be reconstructed from first principles using only the mathematics herein.

Every logical inference has been verified by the machine-checked Coq source corpus comprising the formal verification suite. However, formal verification establishes that *conclusions follow from axioms*—it does not validate that axioms model physical reality. This document is therefore careful to distinguish three epistemological categories: **(S)** Structural theorems about the machine’s own mathematics, **(C)** Conditional derivations where physics conclusions follow from stated axioms that may or may not hold in nature, and **(R)** Consistency relations that verify internal compatibility but do not constitute independent predictions.

The specification covers: the state space and operational semantics; the core “No Free Insight” theorem; computational universality; gauge symmetry and Noether conservation; conditional derivations connecting  $\mu$ -accounting to quantum mechanics (Born rule, Tsirelson bound, no-cloning, unitarity, complex amplitudes, Schrödinger equation) under stated physical axioms; consistency relations (Planck’s constant); emergent spacetime and thermodynamics; the Thiele Manifold and self-reference tower; and falsifiable predictions.

# Contents

<b>1</b>	<b>Introduction</b>	<b>4</b>
1.1	Epistemological Framework . . . . .	4
<b>2</b>	<b>The State Space</b>	<b>6</b>
2.1	Primitives . . . . .	6
2.2	Regions and Modules . . . . .	6
2.3	The Partition Graph . . . . .	6
2.4	Machine State . . . . .	6
<b>3</b>	<b>Instruction Set and Cost Model</b>	<b>7</b>
3.1	Instruction Categories . . . . .	7
3.1.1	Structural Operations . . . . .	7
3.1.2	Logical Operations . . . . .	7
3.1.3	Computational Operations (Reversible ALU) . . . . .	7
3.1.4	Auxiliary Operations . . . . .	7
3.2	Cost Model . . . . .	8
<b>4</b>	<b>Operational Semantics</b>	<b>9</b>
4.1	Transition Function . . . . .	9
4.2	Traces and Execution . . . . .	9
4.3	Selected Instruction Semantics . . . . .	9
<b>5</b>	<b>No Free Insight</b>	<b>10</b>
5.1	Definitions . . . . .	10
5.2	The Theorem . . . . .	10
<b>6</b>	<b>The <math>\mu</math>-Chaitin Theorem</b>	<b>11</b>
<b>7</b>	<b>Cost Equals Kolmogorov Complexity</b>	<b>12</b>
<b>8</b>	<b>No Free Lunch and No Arbitrage</b>	<b>13</b>
8.1	No Free Lunch (Ghosts Are Impossible) . . . . .	13
8.2	No Arbitrage Implies Potential Function . . . . .	13
<b>9</b>	<b>Computational Universality and Subsumption</b>	<b>14</b>
9.1	Turing Machine Embedding . . . . .	14
9.2	Strict Containment . . . . .	14
9.3	Halting Undecidability . . . . .	14
9.4	Oracle Cost Lower Bound . . . . .	14
9.5	Confluence . . . . .	15
9.6	Tensoriality . . . . .	15
9.7	Amortized Analysis . . . . .	15

<b>10</b>	<b><math>\mu</math>-Conservation and Receipt Integrity</b>	<b>16</b>
10.1	The Conservation Law . . . . .	16
10.2	Receipt Integrity . . . . .	16
<b>11</b>	<b>Gauge Symmetry and Noether's Theorem</b>	<b>17</b>
11.1	The $\mathbb{Z}$ -Action (Gauge Group) . . . . .	17
11.2	Gauge Invariance . . . . .	17
<b>12</b>	<b>Two-Dimensional Amplitude Space</b>	<b>18</b>
<b>13</b>	<b>Complex Amplitudes from Norm Preservation</b>	<b>19</b>
<b>14</b>	<b>The Born Rule</b>	<b>20</b>
<b>15</b>	<b>Unitarity</b>	<b>22</b>
<b>16</b>	<b>Purification</b>	<b>23</b>
<b>17</b>	<b>No-Cloning</b>	<b>24</b>
<b>18</b>	<b>Observation and Collapse</b>	<b>25</b>
18.1	Observation Irreversibility . . . . .	25
18.2	Collapse Determination . . . . .	25
<b>19</b>	<b>The Tsirelson Bound</b>	<b>26</b>
<b>20</b>	<b>The Schrödinger Equation</b>	<b>27</b>
<b>21</b>	<b>Planck's Constant</b>	<b>28</b>
<b>22</b>	<b>Emergent Spacetime</b>	<b>29</b>
22.1	The $\mu$ -Metric . . . . .	29
22.2	Causal Cones . . . . .	29
22.3	No-Signaling . . . . .	29
<b>23</b>	<b>Information Causality</b>	<b>30</b>
<b>24</b>	<b>Thermodynamic Bridge</b>	<b>31</b>
24.1	Landauer's Principle . . . . .	31
24.2	Dissipative Embedding . . . . .	31
<b>25</b>	<b>Self-Reference and the Thiele Manifold</b>	<b>32</b>
25.1	Self-Reference Escalation . . . . .	32
25.2	The Thiele Manifold . . . . .	32
25.3	Genesis: Process–Machine Isomorphism . . . . .	32
<b>26</b>	<b>Three-Layer Isomorphism</b>	<b>33</b>
<b>27</b>	<b>The Curry–Howard–Thiele Correspondence</b>	<b>34</b>
<b>28</b>	<b>Categorical Structure</b>	<b>35</b>
<b>29</b>	<b>Audit Infrastructure</b>	<b>36</b>

<b>30 Falsifiable Predictions</b>	<b>37</b>
30.1 Linear Scaling of Structural Cost . . . . .	37
30.2 CHSH Regime Separation . . . . .	37
30.3 Metric Deformation . . . . .	37
30.4 Architectural Permanence . . . . .	37
30.5 Coupling Constant Prediction . . . . .	37
<b>A Proof File Index</b>	<b>39</b>
<b>B Epistemological Classification</b>	<b>41</b>

# Chapter 1

## Introduction

The Thiele Machine is a computational model designed to make the thermodynamic cost of structure explicit. Unlike a Turing Machine, which treats all state transitions as cost-equivalent, the Thiele Machine assigns a specific cost  $\mu$  to every operation that extracts or asserts structural information.

The central claim, proven formally, is: *one cannot gain structural insight about a system without paying  $\mu$ -cost*. This specification explores how this constraint, combined with explicit physical axioms (superposition, linearity, information conservation), yields structural results that parallel the laws of quantum mechanics, thermodynamics, and emergent spacetime. The epistemological status of each result—whether it is a structural theorem (**S**), a conditional derivation (**C**), or a consistency relation (**R**)—is clearly marked throughout.

This specification is organized into eight parts:

1. **Foundations:** State space, instruction set, operational semantics.
2. **Core Theorems:** No Free Insight,  $\mu$ -Chaitin, cost = complexity.
3. **Computation:** Universality, strict subsumption, confluence, halting.
4. **Symmetry & Conservation:**  $\mu$ -conservation, Noether/gauge, receipts.
5. **Quantum Mechanics:** Born rule, Tsirelson, no-cloning, unitarity, complex amplitudes, Schrödinger, Planck.
6. **Emergent Structure:** Spacetime metric, causal cones, information causality, Landauer bridge.
7. **Meta-Theory:** Self-reference, Thiele Manifold, Genesis, three-layer isomorphism, Curry–Howard–Thiele.
8. **Predictions:** Falsifiable divergences from standard QM.

### 1.1 Epistemological Framework

Formal verification (Coq) establishes that conclusions follow logically from axioms. It does **not** establish that axioms model physical reality. This distinction is critical for evaluating the physics claims in Parts V–VIII.

Every result in this specification falls into one of three categories:

- (S) **Structural** Theorems about the Thiele Machine as a mathematical object. These are unconditionally true given the definitions. *Examples:*  $\mu$ -conservation, gauge invariance, confluence, halting undecidability, strict Turing subsumption.

**(C) Conditional** Theorems of the form “if axiom  $X$  holds, then physics result  $Y$  follows.” The logical inference is verified; whether axiom  $X$  holds in nature is an empirical question. *Examples:* Born rule uniqueness (conditional on linearity), complex necessity (conditional on 2D amplitudes), unitarity (conditional on information conservation).

**(R) Consistency** Algebraic identities or definitional equivalences that verify internal coherence but do not constitute independent predictions. *Examples:* Planck consistency relation, Tsirelson cost definition.

Each major theorem in Parts V–VIII is annotated with its category. A theorem marked **(C)** or **(R)** is not diminished—conditional derivations and consistency checks are valuable scientific tools—but the reader should not mistake them for derivations from first principles without stated assumptions. See Appendix B for the complete classification table.

# Chapter 2

## The State Space

### 2.1 Primitives

Let  $\mathbb{N} = \{0, 1, 2, \dots\}$ . Let  $\mathbb{B} = \{0, 1\}$ . Let  $\Sigma$  be a finite alphabet.

### 2.2 Regions and Modules

**Definition 2.1** (Region). A **Region**  $R$  is a finite subset of  $\mathbb{N}$ , canonically represented as a deduplicated list preserving first-occurrence order:  $\text{norm}(L) = \text{nodup}(L)$  removes duplicate elements while preserving the order of their first appearance. (The Python VM applies a stricter canonical form via `sorted(set(region))`; the Coq specification uses `nodup Nat.eq_dec` which does not sort.)

**Definition 2.2** (Module). A **Module**  $M = (R_M, A_M)$  consists of a normalized region  $R_M \subset \mathbb{N}$  and a set of axioms  $A_M \subset \Sigma^*$ .

### 2.3 The Partition Graph

**Definition 2.3** (Partition Graph). A partition graph  $G = (N_{id}, \mathcal{T})$  where  $N_{id} \in \mathbb{N}$  is the next available module ID and  $\mathcal{T} : \{0, \dots, N_{id} - 1\} \rightharpoonup \mathcal{M}$  is a finite partial function from IDs to modules.

**Axiom 2.1** (Well-Formedness).  $G$  is well-formed iff  $\text{dom}(\mathcal{T}) \subseteq \{0, \dots, N_{id} - 1\}$ .

### 2.4 Machine State

**Definition 2.4** (VM State). The complete state is a 7-tuple:

$$S = (G, C, R, \text{Mem}, \text{PC}, \mu, \text{err})$$

where  $G \in \mathcal{G}$  is the partition graph,  $C \in \mathbb{N}^3$  comprises the CSRs (certificate, status, error),  $R \in \mathbb{N}^{32}$  is the register file,  $\text{Mem} \in \mathbb{N}^{256}$  is main memory,  $\text{PC} \in \mathbb{N}$  is the program counter,  $\mu \in \mathbb{N}$  is the  $\mu$ -ledger, and  $\text{err} \in \mathbb{B}$  is the global error flag.

# Chapter 3

# Instruction Set and Cost Model

## 3.1 Instruction Categories

The instruction set  $\mathcal{I}$  is partitioned into three categories. Every instruction  $i$  carries an explicit cost parameter  $\Delta\mu_i \in \mathbb{N}$ .

### 3.1.1 Structural Operations

These modify the partition graph  $G$ :

- $\text{PNEW}(R, \Delta\mu)$ : Create module with region  $R$ .
- $\text{PSPLIT}(m, L, R, \Delta\mu)$ : Split module  $m$  into disjoint regions.
- $\text{PMERGE}(m_1, m_2, \Delta\mu)$ : Merge two modules.
- $\text{PDISCOVER}(m, E, \Delta\mu)$ : Attach evidence  $E$  (list of axioms) to module  $m$ .
- $\text{MDLACC}(m, \Delta\mu)$ : Minimum-description-length accumulate on module  $m$ .

### 3.1.2 Logical Operations

These interact with information content:

- $\text{LASSERT}(m, \phi, cert, \Delta\mu)$ : Assert formula  $\phi$  on module  $m$  with certificate  $cert$ , which is either a SAT model (`lassert_cert_sat`) or an UNSAT LRAT proof (`lassert_cert_unsat`).
- $\text{LJOIN}(c_1, c_2, \Delta\mu)$ : Join certificate checksums.
- $\text{REVEAL}(m, n, cert, \Delta\mu)$ : Reveal  $n$  bits of structure from  $m$ . *Primary source of  $\mu$ -cost*.
- $\text{EMIT}(m, p, \Delta\mu)$ : Emit payload  $p$  from module  $m$ .

### 3.1.3 Computational Operations (Reversible ALU)

Standard reversible register/memory operations: `XFER`, `XOR_LOAD`, `XOR_ADD`, `XOR_SWAP`, `XOR_RANK`.

### 3.1.4 Auxiliary Operations

- $\text{PYEXEC}(payload, \Delta\mu)$ : Execute Python payload (host bridge).
- $\text{CHSH_TRIAL}(x, y, a, b, \Delta\mu)$ : Run a CHSH correlation trial.
- $\text{ORACLE_HALTS}(payload, \Delta\mu)$ : Query halting oracle (requires  $\Delta\mu > 0$ ).

- $\text{HALT}(\Delta\mu)$ : Halt execution.

The full ISA comprises 18 instructions (4 structural + 4 logical + 5 ALU + 5 auxiliary). Every instruction carries an explicit  $\Delta\mu$  parameter; the function `instruction_cost` extracts it.

## 3.2 Cost Model

Every instruction  $i$  carries an explicit cost parameter  $\Delta\mu_i \in \mathbb{N}$ . The function `instruction_cost( $i$ )` extracts this parameter. The cost is applied via  $\text{apply\_cost}(s, i) = s.\mu + \text{instruction\_cost}(i)$ . Since  $\Delta\mu_i : \mathbb{N}$ , the  $\mu$ -ledger is monotonically non-decreasing by construction.

Semantically, only `REVEAL` and `LASSERT` (when adding structure) constitute genuine information-theoretic cost. Other instructions may carry  $\Delta\mu > 0$  as a bookkeeping charge, but in typical programs their cost is set to zero.

# Chapter 4

# Operational Semantics

## 4.1 Transition Function

The transition  $\delta : \mathcal{S} \times \mathcal{I} \rightarrow \mathcal{S}$  is defined by:

1.  $PC' = PC + 1.$
2.  $\mu' = \mu + \Delta\mu_i.$
3.  $err' = err \vee \text{fail}(S, i).$

## 4.2 Traces and Execution

**Definition 4.1** (Trace). A trace  $\tau = [i_0, \dots, i_k] \in \mathcal{I}^*$  is a finite sequence of instructions.

**Definition 4.2** (Execution).  $\text{Run}(\[], S) = S;$   $\text{Run}(i :: \tau', S) = \text{Run}(\tau', \delta(S, i)).$

## 4.3 Selected Instruction Semantics

**PNEW**: Normalize region; if already present, idempotent; else add  $(R_{\text{norm}}, \emptyset)$  to  $G.$

**REVEAL**: Update CSRs with checksum of certificate  $cert.$  Cost:  $\Delta\mu$  (typically 1).

**LASSERT**: Verify certificate  $cert$  for formula  $\phi.$  If  $cert$  is `lassert_cert_unsat`: verify LRAT proof (UNSAT). If  $cert$  is `lassert_cert_sat`: verify model satisfies formula (SAT). If valid, append  $\phi$  to  $A_m;$  if invalid, set  $err \leftarrow \text{true}.$  Replaces oracles with verifiable proofs.

# Chapter 5

## No Free Insight

The central theorem: structural insight is never free.

### 5.1 Definitions

**Definition 5.1** (Receipt Predicate).  $P : \mathcal{O} \rightarrow \mathbb{B}$  is a computable function on observation histories.

**Definition 5.2** (Strength Ordering).  $P_1 \leq P_2$  iff  $\forall o. P_1(o) = 1 \implies P_2(o) = 1$ .  $P_1 < P_2$  iff  $P_1 \leq P_2$  and  $\exists o. P_2(o) = 1 \wedge P_1(o) = 0$ .

**Definition 5.3** (Certification). Trace  $\tau$  **certifies**  $P$  iff execution succeeds ( $err = 0$ ), the supra-certificate flag is set, and the output satisfies  $P$ .

**Definition 5.4** (Structure Addition).  $\text{HasStructureAddition}(\tau)$  is true iff  $\tau$  contains a REVEAL, LASSERT, or equivalent operation.

### 5.2 The Theorem

**Theorem 5.1** (No Free Insight (NoFreeInsight\_Theorem.v)). *Let  $P_{\text{strong}} < P_{\text{weak}}$ . If a trace  $\tau$  starting from a clean state  $s_0$  certifies  $P_{\text{strong}}$ , then:*

$$\text{Certified}(\tau, P_{\text{strong}}) \implies \text{HasStructureAddition}(\tau)$$

*Proof Sketch.* The proof is parametric: `NoFreeInsight_Theorem.v` proves the theorem for any system satisfying the `NO_FREE_INSIGHT_SYSTEM` module type interface. The interface requires four contract obligations (not Coq Axiom declarations—they are proved separately for the kernel instantiation): (1)  $\mu$ -monotonicity, (2) a certification specification linking **certifies** to execution, (3) a no-free-insight contract stating that strict strengthening requires structure addition, and (4) that the system has a notion of clean start. The theorem file itself contains zero `Axiom`, zero `Admitted`. The contract obligations are discharged by the kernel instantiation.  $\square$

**Corollary 5.2** (Cost of Insight). *Gaining insight (strengthening the predicate) implies  $\Delta\mu > 0$ .*

# Chapter 6

## The $\mu$ -Chaitin Theorem

A Chaitin-style incompleteness theorem denominated in  $\mu$ -currency.

**Theorem 6.1** ( $\mu$ -Chaitin (MuChaitinTheory-Theorem.v)). *For any formal theory system  $T$  satisfying the  $\mu$ -Chaitin interface, the number of bits  $k$  the theory can certify is bounded by:*

$$k \leq |T| + c$$

where  $|T|$  is the description length of the theory and  $c$  is a fixed overhead constant. The proof chain:  $k \leq \text{payload} \leq \mu\text{-info} \leq \text{budget} \leq |T| + c$ .

## Chapter 7

# Cost Equals Kolmogorov Complexity

**Theorem 7.1** ( $\mu$ -Bits = Prefix-Free Complexity (CostIsComplexity.v)). Let  $K(\text{spec})$  denote the prefix-free Kolmogorov complexity of specification  $\text{spec}$ . Then:

$$\mu(\text{spec}) = K(\text{spec})$$

Specifically,  $\mu(\text{spec}) = |\text{spec}| + 1$  (the terminating bit), which equals the minimum description length over all prefix-free programs producing  $\text{spec}$ .

*Proof.* The canonical compiler maps  $\text{spec} \mapsto \text{spec} \parallel [\text{true}]$ . Any producing program has the form  $p = \text{spec} \parallel [\text{true}]$  with  $|p| = |\text{spec}| + 1$ . The compiler achieves the minimum, and  $\mu$  counts exactly these bits.  $\square$

# Chapter 8

## No Free Lunch and No Arbitrage

### 8.1 No Free Lunch (Ghosts Are Impossible)

**Theorem 8.1** (No Free Lunch (NoFreeLunch.v)). *A “ghost” is defined as two distinct propositions  $p \neq q$  represented by the same physical state. Ghosts are impossible:*

$$\forall p \neq q, \nexists s \text{ s.t. } s \text{ faithfully represents both } p \text{ and } q$$

*Information cannot exist without physical distinction.*

### 8.2 No Arbitrage Implies Potential Function

**Theorem 8.2** (Potential from No-Arbitrage (NoArbitrage.v)). *Let  $w : \mathcal{I}^* \rightarrow \mathbb{Z}$  be a cost function satisfying:*

1. **Additivity:**  $w([]) = 0$  and  $w(\tau_1 \cdot \tau_2) = w(\tau_1) + w(\tau_2)$ .
2. **No-Arbitrage:** For every closed cycle  $\tau$  (returning to the starting state),  $w(\tau) \geq 0$ .

*Then there exists a potential function  $\phi : \mathcal{S} \rightarrow \mathbb{Z}$  such that:*

$$w(\tau) \geq \phi(\text{apply}(\tau, s)) - \phi(s)$$

*for every state  $s$  and trace  $\tau$ .*

**Remark 8.1.** This is structurally analogous to the Second Law of Thermodynamics: consistent cost accounting implies a potential function bounding all transitions. The concrete Coq model is a toy (increment/decrement on nat with asymmetric costs). No temperature, entropy, or thermodynamic system appears in the formalization. The connection to thermodynamics is an analogy, not a derivation.

# Chapter 9

# Computational Universality and Subsumption

## 9.1 Turing Machine Embedding

**Theorem 9.1** (Abstract Simulation (Embedding\_TM.v, modular\_proofs/TM\_to\_Minsky.v)).  
For every Turing Machine  $T$ , there exists a Minsky counter machine  $M$  and a simulation relation  $R$  such that  $M$  simulates  $T$  step-by-step, encoding the tape into two integers:

$$\text{Left} = \sum_{i=0}^{h-1} t[h-1-i] 2^i, \quad \text{Right} = \sum_{i=0}^{n-h-1} t[h+i] 2^i$$

The Thiele Machine simulates Minsky using its reversible ALU. Hence:

$$\text{TM} \preceq \text{Minsky} \preceq \text{Thiele}$$

## 9.2 Strict Containment

**Theorem 9.2** (Turing  $\subsetneq$  Thiele (Subsumption.v)). Classical Turing computation is **strictly contained** in sighted Thiele computation:

$$\exists p. \text{sighted}(p) \wedge \neg \text{turing}(p)$$

There exist programs that are Thiele-computable (using partition structure and  $\mu$ -accounting) but are not classical Turing programs. The extra structure is genuine new content.

**Theorem 9.3** (Semantic Strictness (ThieleFoundations.v)). There exist Thiele traces with identical Turing shadows but non-isomorphic Thiele-level behavior. The partition and  $\mu$  layers carry information not present in the Turing skeleton.

## 9.3 Halting Undecidability

**Theorem 9.4** (Diagonal Argument (kernel/OracleImpossibility.v)). No total computable function decides the halting problem for the Thiele Machine.

## 9.4 Oracle Cost Lower Bound

**Theorem 9.5** (Oracle Impossibility (kernel/OracleImpossibility.v)). The formalization defines a sound oracle as one charging  $\Delta\mu \geq 1$  per query (`sound_oracle_cost`). Given this definition:

1. *oracle\_halts\_costs\_mu*: A zero-cost oracle violates the soundness definition ( $0 \geq 1 \Rightarrow \perp$ ).
2. *oracle\_cost\_linear*: Given *sound\_oracle\_cost*,  $n$  queries cost  $\geq n$ .

These results verify the internal consistency of the pricing definition. The physical justification for why oracle queries should cost  $\mu > 0$  comes from the No Free Insight framework (gaining information about undecidable propositions should cost something), but the Coq proofs verify the narrower claim that the pricing is self-consistent.

## 9.5 Confluence

**Theorem 9.6** (Church–Rosser (Confluence.v)). *For any state  $s$  and two independent certificates  $c_1, c_2$  (targeting distinct CNF formulas), applying them in either order yields the same  $\mu$ -cost:*

$$\mu(s_{12}) = \mu(s_{21})$$

*Independent structural operations commute.*

## 9.6 Tensoriality

**Theorem 9.7** (Module Independence (ThieleUnificationTensor.v)). *For distinct modules  $m_1 \neq m_2$ , recording discoveries commutes:*

$$\text{record}(\text{record}(G, m_1, e_1), m_2, e_2) \equiv_{\text{ext}} \text{record}(\text{record}(G, m_2, e_2), m_1, e_1)$$

*Local per-module operations are tensorial.*

## 9.7 Amortized Analysis

**Theorem 9.8** (Amortization (AmortizedAnalysis.v)). *For a batch of  $T$  instances with discovery cost  $D$  and operational cost  $O$  per instance:*

$$\text{total\_cost} = B \cdot D + T \cdot O \geq T \cdot O$$

*where  $B$  is batch count. As  $T \rightarrow \infty$ , average cost  $\rightarrow O$  (discovery overhead amortizes away).*

# Chapter 10

## $\mu$ -Conservation and Receipt Integrity

### 10.1 The Conservation Law

**Theorem 10.1** ( $\mu$ -Ledger Conservation (MuLedgerConservation.v)). *For any trace  $\tau = [i_0, \dots, i_k]$ :*

$$\mu_{\text{final}} = \mu_{\text{init}} + \sum_{j=0}^k \text{cost}(i_j)$$

*Every instruction increases  $\mu$  by exactly its declared cost. No axioms, no admits.*

**Corollary 10.2** (Monotonicity).  $\mu$  never decreases:  $\mu(s) \leq \mu(\text{Run}(\tau, s))$  for all  $\tau$ .

**Corollary 10.3** (Irreversibility Bound).  $\text{irreversible\_bits}(\tau) \leq \mu_{\text{final}} - \mu_{\text{init}}$ .

**Theorem 10.4** ( $\mu$  Decomposition (MuLedgerConservation.v)).  $\mu_{\text{total}} = \mu_{\text{blind}} + \mu_{\text{sighted}}$  for any partition of the cost into blind (reversible) and sighted (structural) components.

### 10.2 Receipt Integrity

**Definition 10.1** (Receipt). A receipt  $r = (\text{step}, \text{instr}, \mu_{\text{pre}}, \mu_{\text{post}}, h_{\text{pre}}, h_{\text{post}})$  records one transition.

**Theorem 10.5** (Receipt Chain Validity (ReceiptIntegrity.v)). *A valid receipt chain proves  $\mu_{\text{final}} = \mu_{\text{init}} + \sum \text{costs}$ . Any receipt claiming  $\Delta\mu \neq \text{cost}(\text{instr})$  fails validation.*

**Theorem 10.6** (Non-Forgeability). *Forged receipts (claiming incorrect  $\Delta\mu$ ) fail validation. Overflow values ( $\mu > 2^{31} - 1$ ) are rejected.*

# Chapter 11

## Gauge Symmetry and Noether's Theorem

### 11.1 The $\mathbb{Z}$ -Action (Gauge Group)

**Definition 11.1** (Gauge Shift). For  $\delta \in \mathbb{Z}$ , the gauge shift  $\sigma_\delta : \mathcal{S} \rightarrow \mathcal{S}$  maps  $S \mapsto S[\mu \leftarrow \mu + \delta]$ , preserving all other fields.

**Theorem 11.1** (Group Action (KernelNoether.v)). *The gauge shifts form a  $\mathbb{Z}$ -action on states:*

1. **Identity:**  $\sigma_0(S) = S$ .
2. **Composition:**  $\sigma_a(\sigma_b(S)) = \sigma_{a+b}(S)$ .
3. **Inverse:**  $\sigma_{-n}(\sigma_n(S)) = S$  (when  $\mu + n \geq 0$ ).

### 11.2 Gauge Invariance

**Definition 11.2** (Observable).  $\text{Obs}(S)$  extracts the partition region structure, ignoring  $\mu$ .

**Theorem 11.2** (Gauge Invariance (KernelNoether.v)).  $\text{Obs}(\sigma_\delta(S)) = \text{Obs}(S)$ . *Shifting  $\mu$  by a constant does not affect observables.*

**Theorem 11.3** (Noether's Theorem (KernelNoether.v, RepresentationTheorem.v)).

1. **Forward:** States identical except for  $\mu$  lie on the same gauge orbit. (Trivially true by construction: if all other fields match, the  $\mu$  difference is the gauge shift.)
2. **Gauge Invariance:** `z-gauge_invariance` proves that shifting  $\mu$  does not change `Observable_partition`. This is definitional: `z-gauge_shift` only modifies the `vm_mu` field, and `Observable_partition` reads only `vm_graph`.
3. **Trace Preservation:** Gauge-equivalent states produce identical observable traces (same labels, same  $\mu$ -costs) for any finite horizon. Proven as `vm_step_orbit_equiv`.

The  $\mu$ -ledger is a gauge degree of freedom. Its absolute value is unobservable; only differences  $\Delta\mu$  are physical. The “Noether” label is a deliberate analogy—the result follows from the record structure (separate fields for  $\mu$  and graph), not from a deep symmetry principle.

**Theorem 11.4** ( $\mu$ -Monotonicity (KernelNoether.v)).  $\text{vm\_step}(S, i, S') \implies \mu(S) \leq \mu(S')$ . The  $\mu$ -ledger never decreases under the dynamics. This is true by construction: instruction costs are natural numbers ( $\geq 0$ ), so adding them to the ledger cannot decrease it. The proof examines each step constructor and confirms the arithmetic.

## Chapter 12

# Two-Dimensional Amplitude Space

**Remark 12.1** (Epistemological Status: **(C)** Conditional Derivation). The following derivation assumes that partition states admit *amplitude* representations (superpositions), not merely classical probability distributions. This is stated as Axiom 12.1. Without it, classical probability theory ( $p \in [0, 1]$ , one-dimensional, continuous) suffices and the 2D argument does not apply. The passage from a discrete binary partition to the continuum  $S^1$  also requires a limiting process not formalized in the Coq suite. The argument does not rule out quaternionic (4D) amplitudes; it establishes 2D as the *minimum* dimension for continuous superposition.

**Axiom 12.1** (Superposition Principle). Partition module states admit amplitude representations: a state with  $n$  classical configurations is described by  $n$  real amplitudes  $(a_1, \dots, a_n)$  satisfying  $\sum_i a_i^2 = 1$ , where  $a_i^2$  gives the probability of configuration  $i$ .

Given this axiom, binary partition structure forces quantum amplitudes to live in at least two dimensions.

**Theorem 12.1** (1D Is Insufficient (TwoDimensionalNecessity.v)). *A one-dimensional normalized amplitude satisfies  $x^2 = 1$ , hence  $x \in \{+1, -1\}$ . No intermediate superpositions exist.*

**Theorem 12.2** (2D Is Necessary and Sufficient (TwoDimensionalNecessity.v)). *Binary partition structure requires exactly 2D amplitude space. Two-dimensional amplitudes  $(a, b)$  with  $a^2 + b^2 = 1$  admit a continuous family via  $a = \cos \theta$ ,  $b = \sin \theta$ : the unit circle  $S^1$ .*

## Chapter 13

# Complex Amplitudes from Norm Preservation

Zero-cost evolution must preserve the norm  $a^2 + b^2 = 1$ . The group of norm-preserving maps on  $S^1$  is  $\text{SO}(2) \cong U(1)$ , forcing amplitudes to be complex numbers.

**Theorem 13.1** (Rotation Group (ComplexNecessity.v)). *2D rotations  $R_\theta(a, b) = (a \cos \theta - b \sin \theta, a \sin \theta + b \cos \theta)$  satisfy:*

1.  $|R_\theta(a, b)|^2 = a^2 + b^2$  (norm preservation).
2.  $R_0 = \text{id}$  (identity).
3.  $R_{\theta_1} \circ R_{\theta_2} = R_{\theta_1 + \theta_2}$  (composition = addition).
4.  $R_\theta^{-1} = R_{-\theta}$  (invertibility).

**Theorem 13.2** (Complex Necessity (ComplexNecessity.v)). *Complex multiplication by  $e^{i\theta}$  is exactly 2D rotation. The norm-preserving maps on  $S^1$  are exactly complex multiplications:  $\text{SO}(2) \cong U(1)$ . Hence quantum amplitudes must be complex numbers.*

**Corollary 13.3** (Euler's Formula).  $e^{i\theta_1} \cdot e^{i\theta_2} = e^{i(\theta_1 + \theta_2)}$ .

**Derivation chain:** Binary partition  $\rightarrow$  2D amplitudes  $\rightarrow$  normalization ( $S^1$ )  $\rightarrow$  zero-cost  $\Rightarrow$  norm-preserving  $\Rightarrow \text{SO}(2) \cong U(1) \Rightarrow$  complex numbers.

# Chapter 14

## The Born Rule

**Definition 14.1** (Bloch Sphere Probabilities). For a state with purity vector  $(x, y, z)$ :

$$P(|0\rangle) = \frac{1+z}{2}, \quad P(|1\rangle) = \frac{1-z}{2}$$

**Definition 14.2** (Measurement  $\mu$ -Cost).

$$\text{Cost}(x, y, z) = \frac{1 - (x^2 + y^2 + z^2)}{2}$$

This is the linear entropy. For pure states ( $r^2 = 1$ ), cost = 0. For mixed states ( $r^2 < 1$ ), cost  $> 0$ .

**Theorem 14.1** (Uniqueness of the Born Rule (BornRule.v)). *The Born rule  $P(|0\rangle) = (1+z)/2$  is the **unique** probability assignment satisfying:*

1. Non-negativity and normalization:  $P \geq 0$ ,  $P(|0\rangle) + P(|1\rangle) = 1$ .
2. Eigenstate consistency:  $P(0, 0, 1) = 1$ ,  $P(0, 0, -1) = 0$ .
3. Linearity in  $z$ :  $P(x, y, z, 0) = az + b$  for some  $a, b$ .

*Proof:* Boundary conditions force  $a = 1/2$ ,  $b = 1/2$ . No other solution exists.

**Note:** The Coq theorem also lists a ***mu\_consistent*** hypothesis, but it is **unused** in the proof. The hypothesis ***mu\_consistent***  $P$  unfolds to “measurement cost  $\geq 0$  for valid Bloch vectors,” which is trivially true for all states and does not constrain  $P$ . The linearity assumption (item 3) does all the work—it assumes the functional form  $P = az + b$  and solves a 2-equation system. This is less “deriving the Born rule from  $\mu$ -accounting” and more “the Born rule is the unique linear rule consistent with eigenstates.” The non-circular derivation from tensor consistency (BornRuleFromSymmetry.v, below) is the stronger result.

**Theorem 14.2** (Born Rule from Tensor Consistency (BornRuleFromSymmetry.v)). *Let  $g : [0, 1] \rightarrow [0, 1]$  be any function satisfying:*

1. Eigenstate consistency:  $g(0) = 0$ ,  $g(1) = 1$ .
2. Normalization:  $g(x) + g(1-x) = 1$  for all  $x \in [0, 1]$ .
3. Tensor product consistency:  $g(xy) = g(x) \cdot g(y)$  for  $x, y \in [0, 1]$ .
4. Range:  $0 \leq g(x) \leq 1$  for  $x \in [0, 1]$ .

*Then  $g(x) = x$  for all  $x \in [0, 1]$ .*

**Remark 14.1** (Epistemological Status: **(C)** Conditional Derivation—Circularity Resolved). The original proof (`BornRule.v`) depends on the linearity axiom (item 3):  $P$  is affine in  $z$ . This assumption is equivalent to the Born rule restated, raising a circularity concern.

The new derivation (`BornRuleFromSymmetry.v`) replaces linearity with *tensor product consistency*: independent measurements yield independent outcomes, i.e.,  $g(xy) = g(x)g(y)$ . This is a structural property of the machine, not a physics axiom—it follows from module independence proven in `ThieleUnificationTensor.v`. The proof proceeds:  $g = \text{id}$  on all dyadic rationals (by multiplicativity + normalization); then monotonicity + density of dyadics forces  $g = \text{id}$  everywhere (via the Archimedean property of  $\mathbb{R}$ ). Zero admits.

# Chapter 15

## Unitarity

**Definition 15.1** (Evolution). An evolution  $E$  maps Bloch vectors  $(x, y, z) \mapsto (x', y', z')$  with associated  $\mu$ -cost.

**Theorem 15.1** (Unitarity from Zero Cost (Unitarity.v)).

1.  $\mu = 0 \implies r'^2 \geq r^2$  (*purity cannot decrease at zero cost*). Proven as `zero_cost_preserves_purity`.
2. Non-unitary evolution ( $r'^2 < r^2$  for some state) requires  $\mu > 0$ . Proven as `nonunitary_requires_mu`: the proof specializes the conservation hypothesis at the witness, then `lra`.
3.  $\mu = 0 \Rightarrow$  no info loss (one direction proven). The converse and the full equivalence  $\mu = 0 \Leftrightarrow$  unitary  $\Leftrightarrow$  reversible are stated in the file but not proven as theorems—they appear as `Definition` declarations (propositional constants) rather than proven `Theorems`.

**Theorem 15.2** (Lindblad Requires  $\mu$  (Unitarity.v)). Lindblad-type dissipation at rate  $\gamma > 0$  requires  $\mu \geq \gamma$ . Decoherence is paid for.

**Theorem 15.3** (CPTP Structure (Unitarity.v)). Positivity + trace preservation  $\implies$  the evolution is CPTP (completely positive, trace-preserving).

# Chapter 16

## Purification

**Theorem 16.1** (Purification Principle (Purification.v)). *Every mixed state  $(x, y, z)$  with  $r^2 = x^2 + y^2 + z^2 \leq 1$  admits eigenvalues  $\lambda_1, \lambda_2$  with:*

- $\lambda_1 + \lambda_2 = 1, \quad 0 \leq \lambda_i \leq 1.$
- $(\lambda_1 - \lambda_2)^2 = r^2.$
- *Construction:*  $\lambda_1 = (1 + r)/2, \lambda_2 = (1 - r)/2.$

*Pure states ( $r = 1$ ) need no external reference (deficit = 0). The maximally mixed state ( $r = 0$ ) has deficit = 1.*

# Chapter 17

## No-Cloning

**Theorem 17.1** (No-Cloning from  $\mu$ -Conservation (NoCloning.v)). *Let  $I = x^2 + y^2 + z^2$  be the information content (purity). For a cloning operation with outputs of information  $I_1, I_2$  and  $\mu$ -cost  $\mu$ :*

$$I_1 + I_2 \leq I + \mu$$

**Theorem 17.2** (No-Cloning from  $\mu$ -Monotonicity (NoCloningFromMuMonotonicity.v)). *Machine-native no-cloning using nat-valued structural content (MDL complexity). If cloning duplicates  $n$  bits of structural content while producing zero new bits:  $2n \leq n + 0$  implies  $n \leq 0$ . The entire proof is a single `lia` step (linear integer arithmetic). No Bloch sphere, no Hilbert space, no real analysis. Zero admits.*

**Corollary 17.3** (Perfect Cloning Is Impossible at Zero Cost).  *$I_1 = I_2 = I$  and  $\mu = 0$  implies  $2I \leq I$ , hence  $I \leq 0$ . Only the trivial state can be cloned for free.*

**Corollary 17.4** (Cloning Cost). *Perfect cloning requires  $\mu \geq I$ . For pure states:  $\mu \geq 1$ .*

**Theorem 17.5** (Approximate Cloning (NoCloning.v)). *For fidelities  $f_1, f_2$ :  $f_1 + f_2 \leq 1 + \mu/I$ . At zero cost:  $f_1 + f_2 \leq 1$ .*

**Theorem 17.6** (No Deletion (NoCloning.v)). *Perfect deletion also requires  $\mu \geq I$  (dual of no-cloning).*

# Chapter 18

## Observation and Collapse

### 18.1 Observation Irreversibility

**Theorem 18.1** (REVEAL Is Irreversible (ObservationIrreversibility.v)). *For bits > 0:  $\mu_{\text{after}} > \mu_{\text{before}}$ , and the post-REVEAL state  $\neq$  the pre-REVEAL state. Observation prevents recovery of the pre-measurement superposition.*

### 18.2 Collapse Determination

**Definition 18.1** (Partition Entropy).  $H(P) = \log_2(\dim(P))$  where  $\dim$  is the partition state dimension.

**Theorem 18.2** (Maximum Information Implies Unique State (CollapseDetermination.v)). *If bits revealed =  $H(P_{\text{before}})$  (maximum information), then  $\dim(P_{\text{after}}) = 1$ . The measurement fully collapses the state to a unique outcome. This is the **projection postulate**, derived rather than assumed.*

# Chapter 19

## The Tsirelson Bound

**Definition 19.1** (Correlation  $\mu$ -Cost).  $\mu_{\text{corr}}(S) = 0$  if  $|S| \leq 2\sqrt{2}$ ; otherwise  $\mu_{\text{corr}} = |S| - 2\sqrt{2}$ .

**Remark 19.1** (Epistemological Status: **(C)** Conditional Derivation—Circularity Resolved). The original Coq proof (`TsirelsonDerivation.v`) encoded  $2\sqrt{2}$  in the definition of  $\mu_{\text{corr}}$ , making its derivation status **(R)**.

This gap is now closed. `TsirelsonGeneral.v` derives  $S^2 \leq 8$  from pure algebra: a sum-of-squares identity proves each CHSH row contributes  $\leq 2(e_1^2 + e_2^2)$ , then Cauchy–Schwarz bounds the sum by  $4 \sum e_i^2 \leq 8$ . `TsirelsonFromAlgebra.v` provides a self-contained algebraic derivation with achievability witness ( $e = \pm 1/\sqrt{2}$ , giving  $S = 2\sqrt{2}$ ) and a verified rational bound ( $\sqrt{8} < 5657/2000$ ). Zero admits in both files.

**Theorem 19.1** (Tsirelson from Zero  $\mu$  (`TsirelsonDerivation.v`) — **(R)**). *Total  $\mu = 0$  implies  $|S_{\text{CHSH}}| \leq 2\sqrt{2} \approx 2.828$ .*

**Theorem 19.2** (CHSH Separation (`Deliverable_CHSHSeparation.v`)). *Strict numerical separation:*

$$2 \quad < \quad 2\sqrt{2} \approx 2.828 \quad < \quad \frac{16}{5} = 3.2$$

- **Classical** (local receipts):  $|S| \leq 2$  (Bell bound).
- **Quantum** (admissible,  $\mu = 0$ ):  $|S| \leq 2\sqrt{2}$  (Tsirelson bound).
- **Supra-quantum** ( $\mu > 0$  allowed): witnesses achieve  $|S| = 3.2$ .

*The Tsirelson bound  $2\sqrt{2}$  is the maximum correlation purchasable at zero  $\mu$ -cost. Exceeding it requires paying for additional structural information.*

## Chapter 20

# The Schrödinger Equation

**Theorem 20.1** (Emergent Schrödinger (EmergentSchrodinger.v)). *The finite-difference Schrödinger equation emerges from the partition update rules. For a two-component amplitude state  $(a, b)$  with mass  $m$ , potential  $V$ , and time step  $\Delta t$ :*

$$a(t + \Delta t) = a - \frac{\Delta t}{2m} \nabla^2 b + \Delta t \cdot V b \quad (20.1)$$

$$b(t + \Delta t) = b + \frac{\Delta t}{2m} \nabla^2 a - \Delta t \cdot V a \quad (20.2)$$

The coupling is **antisymmetric**:  $c_{a,\nabla^2 b} = -c_{b,\nabla^2 a}$  and  $c_{a,Vb} = -c_{b,Va}$ , which is necessary and sufficient for probability conservation  $\partial_t(a^2 + b^2) = 0$ .

**Remark 20.1** (Epistemological Status: (C) Conditional Derivation). This is exactly the finite-difference discretization of  $i\hbar\partial_t\psi = -\frac{\hbar^2}{2m}\nabla^2\psi + V\psi$  with  $\psi = a + ib$ . The Schrödinger form emerges *conditionally*: given Axiom 12.1 (which entails 2D amplitudes) and the requirement that zero-cost evolution preserves probability ( $\partial_t(a^2 + b^2) = 0$ ), the antisymmetric coupling structure is the unique linear solution. The mass  $m$  and potential  $V$  are parameters, not derived.

# Chapter 21

## Planck's Constant

**Remark 21.1** (Epistemological Status: **(R)** Consistency Relation). The following is **not** a derivation of Planck's constant from first principles. It is a consistency relation that gives  $h$  a physical interpretation within the  $\mu$ -framework. The Coq proof verifies an algebraic identity; the physics is in the definitions.

**Definition 21.1** (Physical Constants (Normalized Units)).  $k_B = 1/100$ ,  $T = 1$ ,  $h = 1$ . Landauer energy:  $E_L := k_B T \ln 2$ .

**Definition 21.2** (Computational Time Step).  $\tau_\mu := h/(4E_L)$ , the Margolus–Levitin time at Landauer energy.

**Theorem 21.1** (Planck Consistency Relation (PlanckDerivation.v) — **(R)**).

$$h = 4 \cdot E_L \cdot \tau_\mu$$

**Remark 21.2** (Why This Is Circular). Substituting  $\tau_\mu = h/(4E_L)$  yields  $h = 4E_L \cdot h/(4E_L) = h$ . The Coq proof reduces to the `field` tactic (pure algebra). This does *not* predict  $h$ ; it defines  $\tau_\mu$  in terms of  $h$  and verifies consistency.

**What the relation does provide:** a physical interpretation of Planck's constant as  $4 \times$  the action (energy  $\times$  time) of one  $\mu$ -bit operation at Landauer cost. If taken as a physical claim, the testable content is the predicted time step  $\tau_\mu = h/(4k_B T \ln 2)$ , which at room temperature yields  $\tau_\mu \approx 5.7 \times 10^{-14}$  s (femtosecond range, consistent with molecular vibration timescales).

**What would make this non-circular:** defining  $\tau_\mu$  as an independent primitive (measured experimentally) and then *deriving*  $h = 4E_L\tau_\mu$  as a *prediction*. The current formalization does not do this.

## Chapter 22

# Emergent Spacetime

### 22.1 The $\mu$ -Metric

**Definition 22.1** ( $\mu$ -Distance (MuGeometry.v)). The Coq formalization defines `mu_distance` as the absolute difference of  $\mu$ -totals:

$$d_\mu(S_1, S_2) = |\mu_{\text{total}}(S_2) - \mu_{\text{total}}(S_1)|$$

This is a one-dimensional projection, not a minimum over paths.

**Theorem 22.1** (Pseudometric Properties (MuGeometry.v)).  $d_\mu$  satisfies non-negativity, reflexivity ( $d(S, S) = 0$ ), symmetry, and triangle inequality—all inherited from absolute value on  $\mathbb{Z}$ . Note:  $d_\mu$  is a **pseudometric**, not a metric:  $d(S_1, S_2) = 0$  does not imply  $S_1 = S_2$  (many distinct states share the same  $\mu$ -total). The “geometry” is a 1D projection onto the  $\mu$  counter.

### 22.2 Causal Cones

**Definition 22.2** ( $\mu$ -Reachability Cone (SpacetimeEmergence.v)).  $C^+(S) = \{S' \mid \exists \tau. S \xrightarrow{\tau} S' \wedge \text{cost}(\tau) \leq \text{Budget}\}$ .

This is a structural analogy to causal cones: the set of reachable states given a finite  $\mu$ -budget plays a role analogous to a light cone. The analogy should not be over-read—there is no Lorentz invariance, no metric signature, and no boost transformations in the Coq formalization.

### 22.3 No-Signaling

**Theorem 22.2** (Observational Locality (SpacetimeEmergence.v)). If module  $M_B$  is not in the target set of any instruction in a trace, then  $M_B$ ’s observable region is unchanged after execution (`exec_trace_no_signaling_outside_cone`).

**Remark 22.1.** This is a **data isolation** property: operations targeting module  $A$  in the partition graph do not affect the lookup of module  $B$ . The proof is structural (list operations on disjoint keys), non-trivial mainly for `psplit/pmerge` which restructure the graph. The label “no-signaling” is an analogy—there is no spatial metric or speed limit between modules.

## Chapter 23

# Information Causality

**Theorem 23.1** (IC– $\mu$  Equivalence (InformationCausality.v) — (R)). *The file defines two record types (`ICScenario` and `MuScenario`) with opaque `Prop` fields, then proves they are equivalent given a hypothesis (`ic_mu_equivalent`) that literally states the equivalence. The proof is tauto. No mutual information, probability distributions, or CHSH values appear anywhere in the formalization.*

**Remark 23.1.** The intended statement—that Bob’s total information about Alice’s data is bounded by the channel’s  $\mu$ -capacity—is a reasonable physical claim but is **not formalized** in the Coq file. The current file is a placeholder that assumes the conclusion as a hypothesis and re-derives it. A future version would need to define entropy, mutual information, and prove the bound from the VM semantics.

# Chapter 24

## Thermodynamic Bridge

### 24.1 Landauer's Principle

**Theorem 24.1** (Landauer Derived (LandauerDerived.v)).

1. *erasure\_irreversible*: Erasing  $\geq 1$  bit implies fan-in =  $2^k > 1$  (arithmetic:  $2^k > 1$  for  $k \geq 1$ ).
2. *erasure\_decreases\_entropy*: Erasure decreases system entropy ( $\Delta S < 0$ ) — this is just output < input  $\Rightarrow$  output – input < 0.
3. *landauer\_information\_bound*: For any physical erasure, environment entropy increase  $\geq$  bits erased.
4. *erasure\_additive*: Sequential erasures compose additively.

*Caveat on item 3: The PhysicalErasure record has second\_law\_satisfied as a field, not a derived property. The theorem landauer\_information\_bound simply extracts this field, making the bound an input assumption rather than a derivation. The physical energy bound  $Q \geq k_B T \ln 2 \cdot n$  is explicitly not proven (the file itself notes the conversion factor is not formalized). Also, the mu in this file is a standalone synonym for bits\_erased—it has no formal connection to vm\_mu in VMState.v.*

### 24.2 Dissipative Embedding

**Theorem 24.2** (Dissipative Model (DissipativeEmbedding.v)). *Irreversible (dissipative) physics embeds into the Thiele VM with  $\Delta\mu \geq 1$  per irreversible step. Reversible physics embeds with  $\Delta\mu = 0$ .*

**Theorem 24.3** (Reversible Physics (PhysicsEmbedding.v, WaveEmbedding.v)). *Reversible lattice gas and wave models embed with  $\mu_{\text{final}} = \mu_{\text{init}}$  (zero cost). Particle count, momentum, wave energy, and wave momentum are conserved by the VM.*

# Chapter 25

## Self-Reference and the Thiele Manifold

### 25.1 Self-Reference Escalation

**Theorem 25.1** (Gödel-Style Escalation (SelfReference.v)). *Any self-referential system  $S$  requires a meta-system  $\text{Meta}$  with strictly more dimensions:*

$$\text{self\_ref}(S) \implies \exists \text{Meta} : \dim(\text{Meta}) > \dim(S)$$

*The meta-system inherits self-reference, creating an infinite escalation.*

### 25.2 The Thiele Manifold

**Definition 25.1** (Thiele Manifold (ThieleManifold.v)). *An infinite tower of systems  $\{L_n\}_{n \in \mathbb{N}}$  with:*

- $\dim(L_n) = 4 + n$  (level  $n$  has dimension  $4 + n$ ).
- $\dim(L_{n+1}) > \dim(L_n)$  (strict enrichment).
- Each level reasons about the level below.
- Self-reference at level  $n$  is answered by level  $n + 1$ .

**Theorem 25.2** (Projection to Spacetime (ThieleManifold.v)). *The projection  $\pi_4$  collapses the tower to dimension 4 (spacetime). For  $n > 0$ :*

- $\pi_4$  is lossy:  $\dim(L_n) > 4$ .
- $\mu$ -cost of projection =  $\dim(L_n) - 4 = n > 0$ .

*Spacetime is the “shadow” of the full manifold.*

### 25.3 Genesis: Process–Machine Isomorphism

**Theorem 25.3** (Genesis (Genesis.v)). *There is a definitional isomorphism  $\text{Proc} \cong \text{Thiele}$  between coherent processes (step + admissibility proof) and Thiele machines:*

$$\text{thiele\_to\_proc}(\text{proc\_to\_thiele}(P)) = P \tag{25.1}$$

$$\text{proc\_to\_thiele}(\text{thiele\_to\_proc}(T, H)) = T \tag{25.2}$$

*Any coherent process is a Thiele machine, and vice versa.*

## Chapter 26

# Three-Layer Isomorphism

**Theorem 26.1** (Full Isomorphism (FullIsomorphism.v)). *Three implementations — Coq specification, Python VM, and Verilog RTL — are isomorphic: for any trace  $\tau$ :*

$$\text{decode}(S_{\text{Coq}}(\tau)) = \text{decode}(S_{\text{Python}}(\tau)) = \text{decode}(S_{\text{Verilog}}(\tau))$$

*with  $\mu(\text{run}(s, \tau)) = \mu(s) + \sum_i \text{cost}(\tau_i)$  at every layer. Transitivity:  $\text{Coq} \cong \text{Python} \cong \text{Verilog} \implies \text{Coq} \cong \text{Verilog}$ .*

## Chapter 27

# The Curry–Howard–Thiele Correspondence

**Theorem 27.1** (Logic–Computation Isomorphism (LogicIsomorphism.v)). *The Thiele Machine extends the Curry–Howard correspondence:*

<i>Logic</i>	<i>Computation</i>	<i>Thiele</i>
<i>Proposition</i>	<i>Type</i>	<i>Partition</i>
<i>Proof</i>	<i>Program (term)</i>	<i>Trace</i>
<i>Cut elimination</i>	$\beta$ - <i>reduction</i>	<i>Execution (Run)</i>
<i>Proof equivalence</i>	$=_\beta$	<i>Execution equivalence</i>

*A valid proof corresponds to a terminating execution. Proof equivalence  $\Leftrightarrow$  execution equivalence.*

**Theorem 27.2** (Logic Is Physics (LogicToPhysics.v)). *Cut elimination in logic = relational composition in physics:*

$$\text{interp}(\text{cut}(\pi_1, \pi_2)) = \text{rel\_comp}(\text{interp}(\pi_1), \text{interp}(\pi_2))$$

*Logical proof composition maps to physical relation composition. This is the formal nucleus of the claim that logic and physics share the same categorical structure.*

## Chapter 28

# Categorical Structure

**Theorem 28.1** (Conservation as Functor (Universe.v)). *Define two categories:*

- $\mathbf{C}_{\text{phys}}$ : Objects = universe states (particle momenta lists). Morphisms = paths of momentum-conserving interactions.
- $\mathbf{C}_{\text{logic}}$ : Objects = total momentum values. Morphisms = equality proofs.

The functor  $F : \mathbf{C}_{\text{phys}} \rightarrow \mathbf{C}_{\text{logic}}$  maps  $F(s) = \sum s$  (total momentum). Then:

$$\text{Path}(s_1, s_2) \implies \sum s_1 = \sum s_2$$

Conservation of momentum emerges as the functorial image. Observation is a structure-preserving map from physics to logic.

## Chapter 29

# Audit Infrastructure

**Theorem 29.1** (CatNet Integrity (CatNet.v)). *Adding a new entry to a valid hash-chain audit log produces a valid chain. If the logic oracle detects inconsistency, execution halts and no further entries are written (paradox halting).*

# Chapter 30

## Falsifiable Predictions

### 30.1 Linear Scaling of Structural Cost

**Axiom 30.1** (Linear Scaling (kernel/FalsifiablePrediction.v)). The  $\mu$ -cost of maintaining coherence of  $N$  entangled qubits scales as  $O(N)$  per time step.

**Prediction:** Large-scale quantum computers will encounter a fundamental (not merely technical) decoherence noise floor proportional to entanglement complexity.

### 30.2 CHSH Regime Separation

**Prediction:** Experiments probing the boundary between quantum ( $|S| \leq 2\sqrt{2}$ ) and supra-quantum ( $|S| > 2\sqrt{2}$ ) correlations should find that exceeding the Tsirelson bound requires measurably higher energy dissipation, scaling as:

$$\Delta E \geq k_B T \ln 2 \cdot (|S| - 2\sqrt{2})$$

### 30.3 Metric Deformation

Since  $d_\mu$  depends on information content, regions of high structural complexity effectively expand the metric.

**Prediction:** In high-complexity computations, effective signal latency will increase relative to vacuum propagation. (This prediction relies on interpreting  $d_\mu$  as physically meaningful, which is currently an analogy rather than a derived result.)

### 30.4 Architectural Permanence

**Theorem 30.1** (Optimal Quartet (ArchTheorem.v)). *The four partition discovery strategies (Louvain, Spectral, Degree, Balanced) achieve classification accuracy > 90% across all problem classes. No alternative configuration exceeds this quartet's accuracy. The configuration is architecturally final.*

### 30.5 Coupling Constant Prediction

**Definition 30.1** (Thiele  $\alpha$  Limit (thiele\_manifold/PhysicalConstants.v)). The asymptotic density of self-referential programs in the  $n$ -bit state space:

$$\alpha = \lim_{n \rightarrow \infty} \frac{A_{\text{interaction}}(n)}{V_{\text{spacetime}}(n)} = \lim_{n \rightarrow \infty} \frac{n+1}{2^n}$$

This converges to 0, but the *non-asymptotic* structure at physical scales should relate to coupling constants.

## **Appendix A**

## **Proof File Index**

The following table maps each theorem in this specification to its Coq source file.

Result	Coq Source
No Free Insight (Thm. 5.1)	nofi/NoFreeInsight_Theorem.v
$\mu$ -Chaitin	nofi/MuChaitinTheory_Theorem.v
Cost = Complexity	theory/CostIsComplexity.v
No Free Lunch	theory/NoFreeLunch.v
No Arbitrage	kernel/NoArbitrage.v
TM Embedding	thielemachine/coqproofs/Embedding_TM.v
TM $\rightarrow$ Minsky	modular_proofs/TM_to_Minsky.v
Strict Subsumption	thielemachine/coqproofs/Subsumption.v
Semantic Strictness	thielemachine/coqproofs/ThieleFoundations.v
Confluence	thielemachine/coqproofs/Confluence.v
Tensoriality	thielemachine/coqproofs/ThieleUnificationTensor.v
$\mu$ -Conservation	kernel/MuLedgerConservation.v
Receipt Integrity	kernel/ReceiptIntegrity.v
Noether / Gauge	kernel/KernelNoether.v
Gauge Trace Preservation	thielemachine/coqproofs/RepresentationTheorem.v
2D Necessity	quantum_derivation/TwoDimensionalNecessity.v
Complex Necessity	quantum_derivation/ComplexNecessity.v
Born Rule	kernel/BornRule.v
Born Rule (Tensor)	kernel/BornRuleFromSymmetry.v
Unitarity	kernel/Unitarity.v
Purification	kernel/Purification.v
No-Cloning	kernel/NoCloning.v
No-Cloning (Machine-Native)	kernel/NoCloningFromMuMonotonicity.v
Observation Irreversibility	quantum_derivation/ObservationIrreversibility.v
Collapse Determination	quantum_derivation/CollapseDetermination.v
Tsirelson Derivation	kernel/TsirelsonDerivation.v
Tsirelson (Algebraic)	kernel/TsirelsonFromAlgebra.v
Tsirelson (General)	kernel/TsirelsonGeneral.v
CHSH Separation	thielemachine/verification/Deliverable_CHSHSeparation.v
Emergent Schrödinger	physics_exploration/EmergentSchrodinger.v
Planck Consistency	physics_exploration/PlanckDerivation.v
Spacetime Emergence	kernel/SpacetimeEmergence.v
$\mu$ -Geometry	kernel/MuGeometry.v
Information Causality	kernel/InformationCausality.v
Landauer	thermodynamic/LandauerDerived.v
Dissipative Embedding	thielemachine/coqproofs/DissipativeEmbedding.v
Physics Embedding	thielemachine/coqproofs/PhysicsEmbedding.v
Wave Embedding	thielemachine/coqproofs/WaveEmbedding.v
Self-Reference	self_reference/SelfReference.v
Thiele Manifold	thiele_manifold/ThieleManifold.v
Genesis	theory/Genesis.v
Full Isomorphism	thielemachine/verification/FullIsomorphism.v
Curry–Howard–Thiele	thielemachine/coqproofs/LogicIsomorphism.v
Logic Is Physics	theory/LogicToPhysics.v
Functor Soundness	isomorphism/coqproofs/Universe.v
CatNet Integrity	catnet/coqproofs/CatNet.v
Amortization	thielemachine/coqproofs/AmortizedAnalysis.v
Optimal Quartet	theory/ArchTheorem.v
Falsifiable Prediction	kernel/FalsifiablePrediction.v
Physical Constants	thiele_manifold/PhysicalConstants.v
Halting Undecidability	kernel/OracleImpossibility.v
Oracle Cost Bound	kernel/OracleImpossibility.v
Hyper-Thiele Oracle	thielemachine/coqproofs/HyperThiele_Halting.v

## Appendix B

# Epistemological Classification

The following table classifies every major result by its epistemological status.

Result	Status	Key Assumption
<i>Unconditional structural theorems about the machine</i>		
No Free Insight	(S)	Module type contract
$\mu$ -Chaitin	(S)	Module type contract
Cost = Complexity	(S)	Prefix-free coding
No Free Lunch	(S)	Faithful representation
No Arbitrage $\Rightarrow$ Potential	(S)	Additivity + no-arbitrage (toy model)
TM Embedding	(S)	Standard simulation
Strict Subsumption	(S)	Partition structure
Halting Undecidability	(S)	Diagonalization
Oracle Cost	(S)	Definitional (soundness := cost $\geq 1$ )
Confluence	(S)	Module independence
$\mu$ -Conservation	(S)	Machine semantics only
Receipt Integrity	(S)	Hash-chain model
Gauge / Noether	(S)	Definitional (record structure)
Genesis	(S)	Definitional isomorphism
Three-Layer Isomorphism	(S)	Decode function
Curry–Howard–Thiele	(S)	Interpretation map
<i>Conditional derivations (valid given stated axioms)</i>		
2D Necessity	(C)	Superposition axiom (Axiom 12.1)
Complex Necessity	(C)	2D + norm preservation
Born Rule Uniqueness	(C)	Tensor product consistency ( <code>BornRuleFromSymmetry.v</code> )
Unitarity from $\mu = 0$	(C)	Info conservation (one direction only)
No-Cloning	(C)	Purity = information
Purification	(C)	Bloch sphere model
Observation Irreversibility	(C)	REVEAL semantics
Collapse Determination	(C)	Entropy = log dim
Emergent Schrödinger	(C)	2D + antisymmetry
Landauer	(C)	Second law
emphassumed as record field		
Dissipative Embedding	(C)	Irreversibility model
Self-Reference	(C)	Dimension = complexity
CHSH Separation	(C)	Numerical witnesses
Tsirelson Bound	(C)	Algebraic derivation ( <code>TsirelsonFromAlgebra.v</code> )
<i>Consistency relations (definitions verified, not predictions)</i>		
Planck's Constant	(R)	$\tau_\mu := h/(4E_L)$
Information Causality	(R)	Circular: equivalence assumed as hypothesis
$\mu$ -Geometry	(R)	Pseudometric (absolute value on $\mathbb{Z}$ )
Spacetime Locality	(R)	Data isolation in partition graph
Born Rule ( <code>BornRule.v</code> )	(R)	$\mu$ -consistent hypothesis unused

**Legend:** (S) = unconditional theorem about the machine; (C) = valid derivation conditional

on stated axiom; **(R)** = internal consistency check.